

CONFIDENTIAL: Oxmedica Ltd. 10347756



Day 8



CONFIDENTIAL: Oxmedica Ltd. 10347756



Overview

1

Visualisations

2

Linear regression as a computational graph

3

An introduction to neural networks

4

Assessment feedback

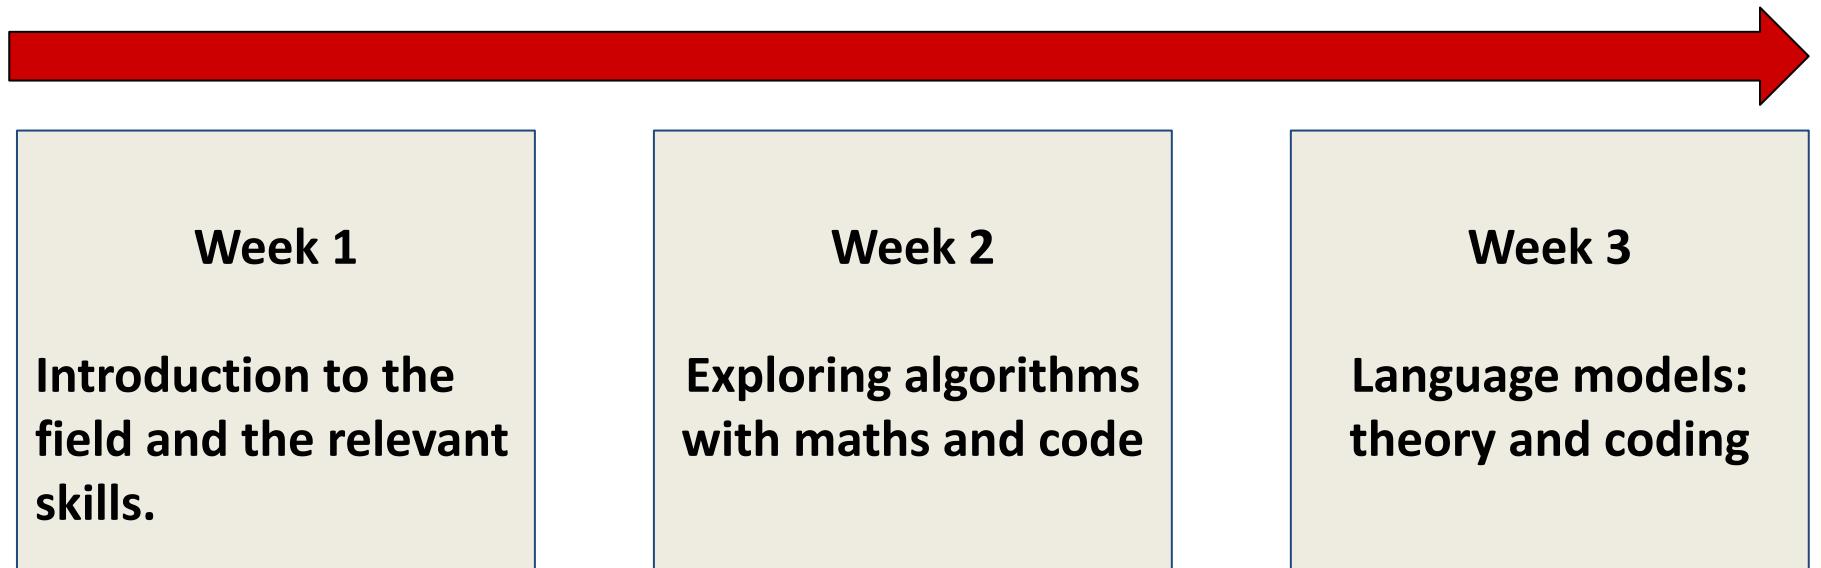


CONFIDENTIAL: Oxmedica Ltd. 10347756

Halfway point

30th June

18th July



Name	Points
Ahmed K	24
Naser	20
Turki	18

Visualisations



Best Visualisations



موهبة

مؤسسة الملك عبد العزيز ورجاله للموهبة والإبداع
King Abdulaziz & Sons Research Foundation for Genetics & Society

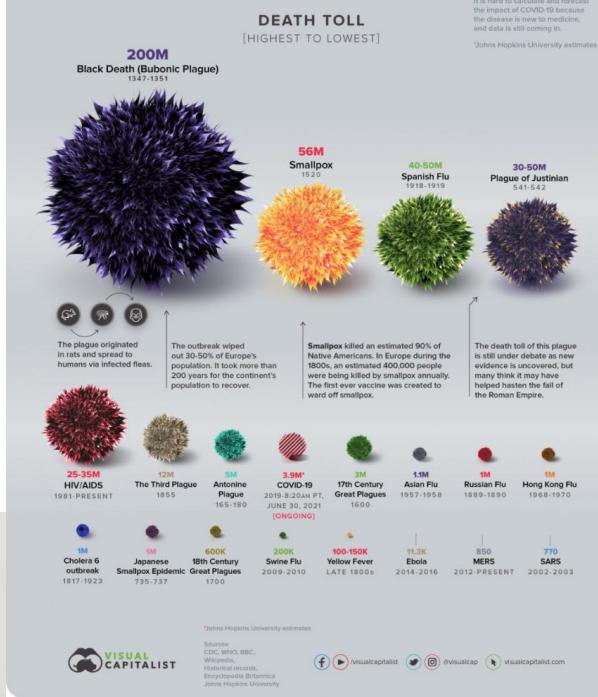
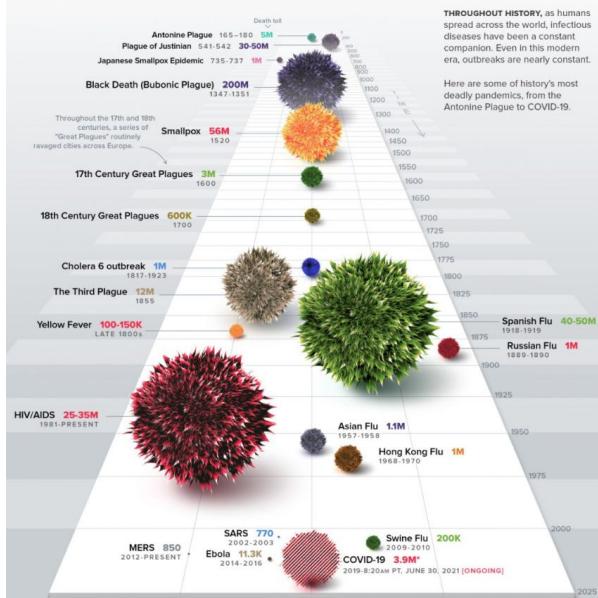
CONFIDENTIAL: Oxmedica Ltd. 10347756



1

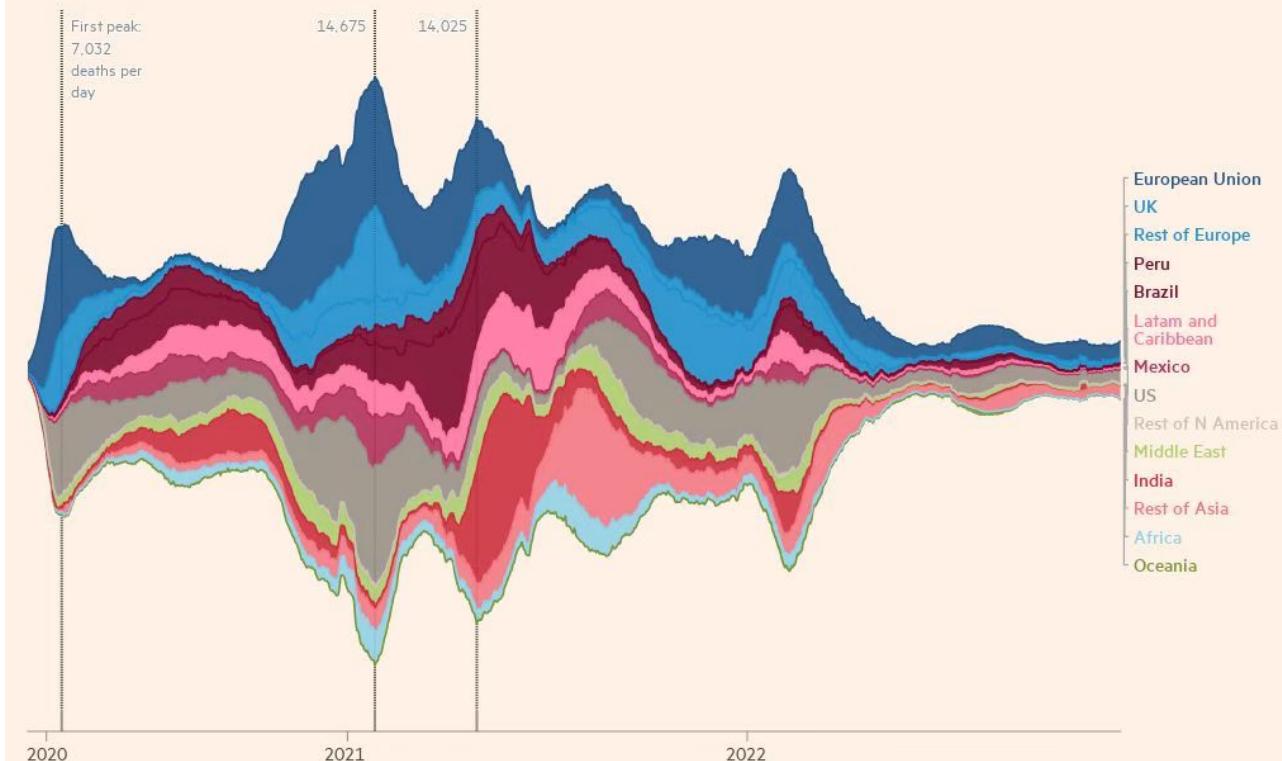
HISTORY OF PANDEMICS

PAN-DEM-IC (of a disease) prevalent over
a whole country or the world.



More than 1,000 deaths each day are attributed to Covid-19

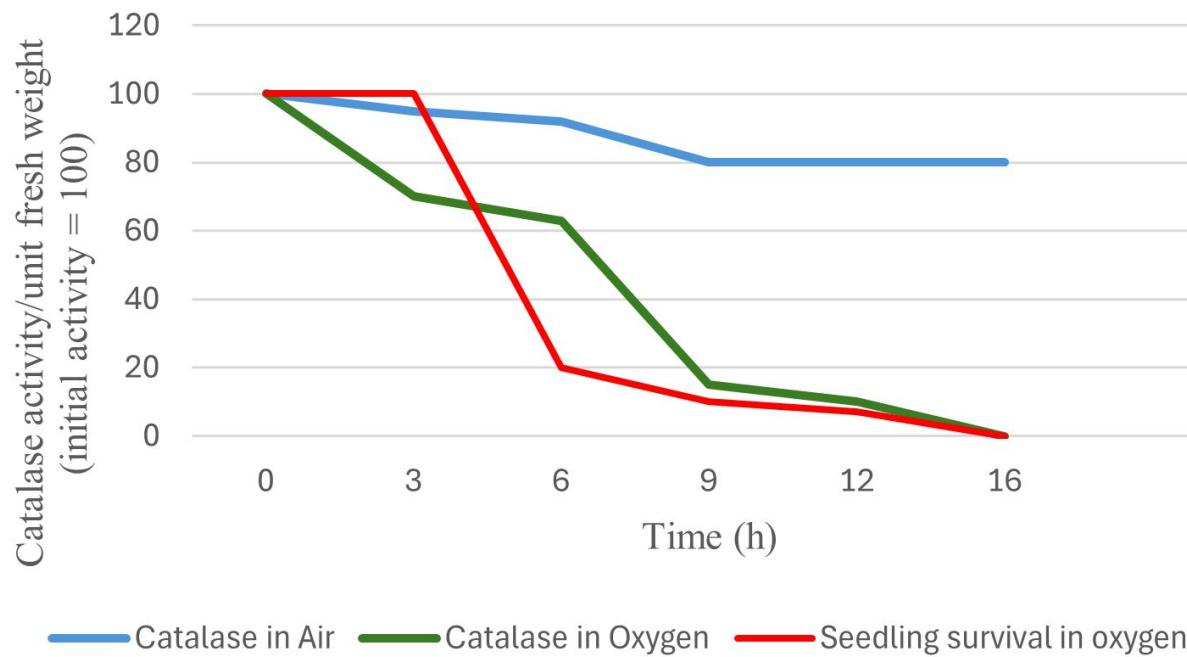
Daily deaths attributed to Covid-19 (7-day rolling average)



Source: Johns Hopkins CSSE, WHO, national sources, FT research • Excludes recent data covering less than 95% of global population
N America includes Canada, Bermuda, Greenland and St Pierre and Miquelon

FINANCIAL TIMES

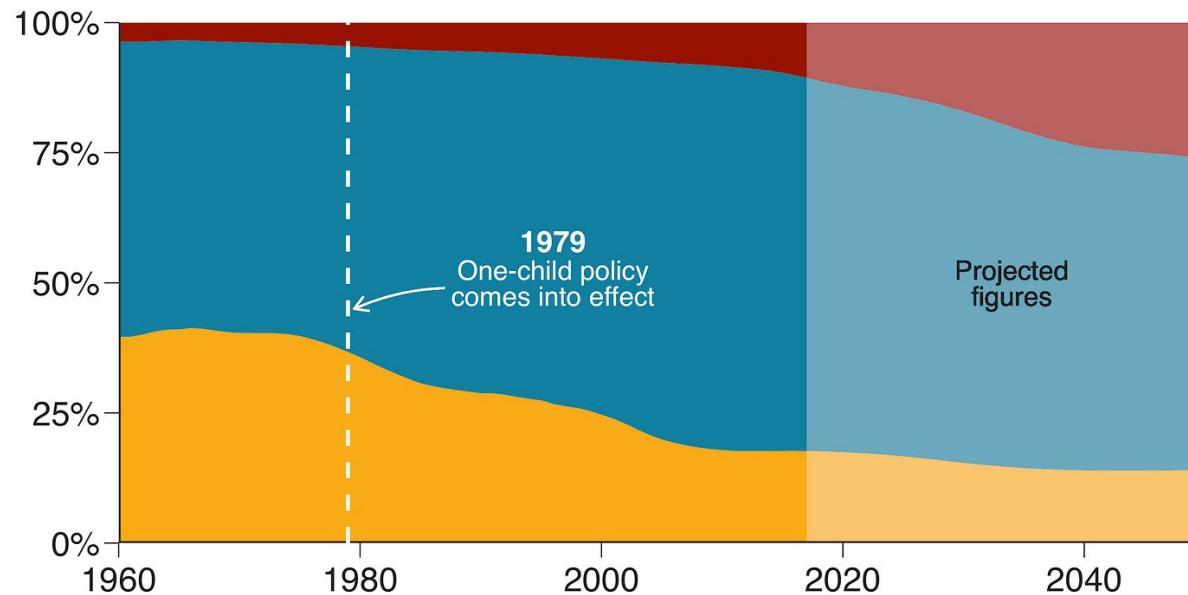
Figure 1. Effect of treatment with 10 atm. pure oxygen on catalase activity and seedling mortality of cress seedlings.



Breakdown of China's population by age group

Proportion of total population (1960-2050)

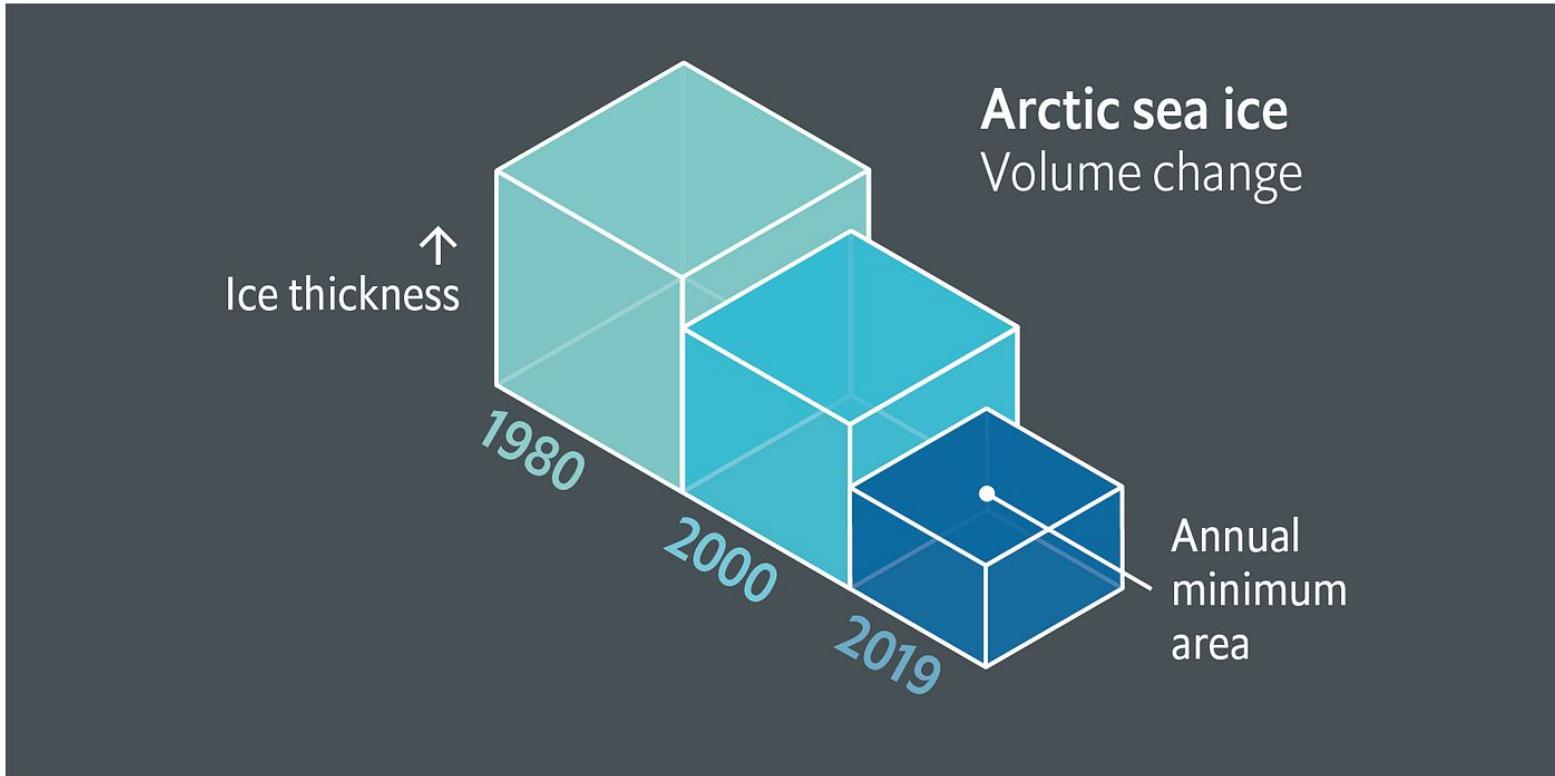
■ 0-14 years ■ 15-64 ■ 65+



Source: The World Bank

BBC

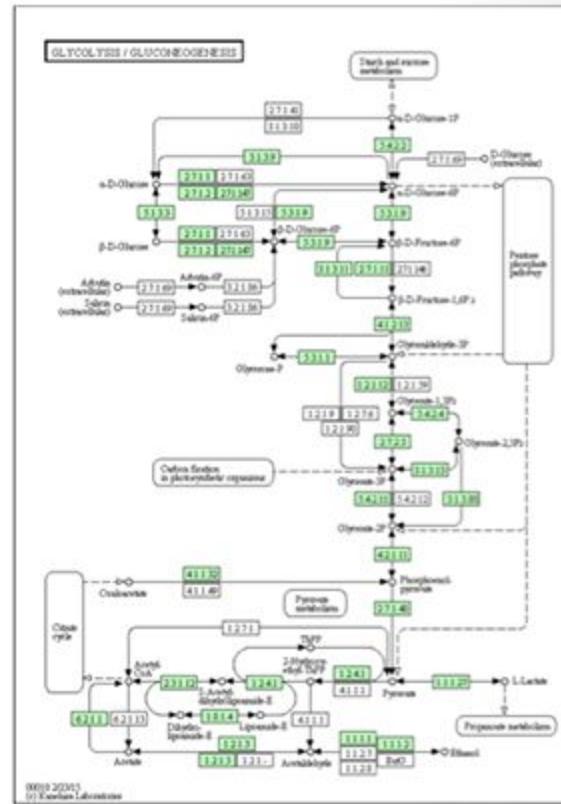
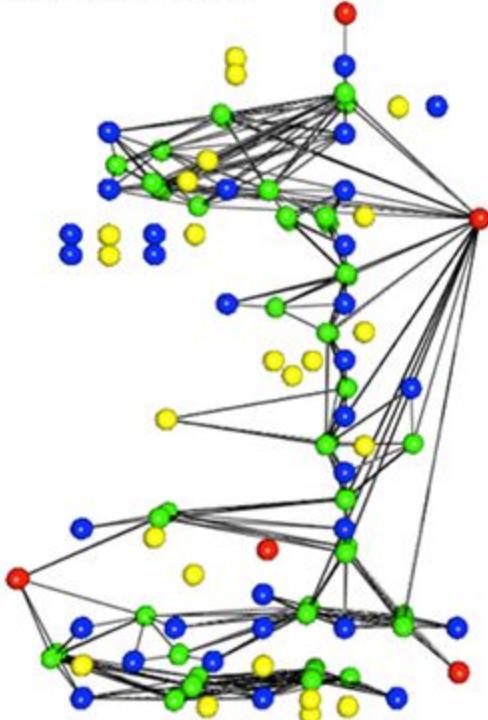




Worst Visualisations



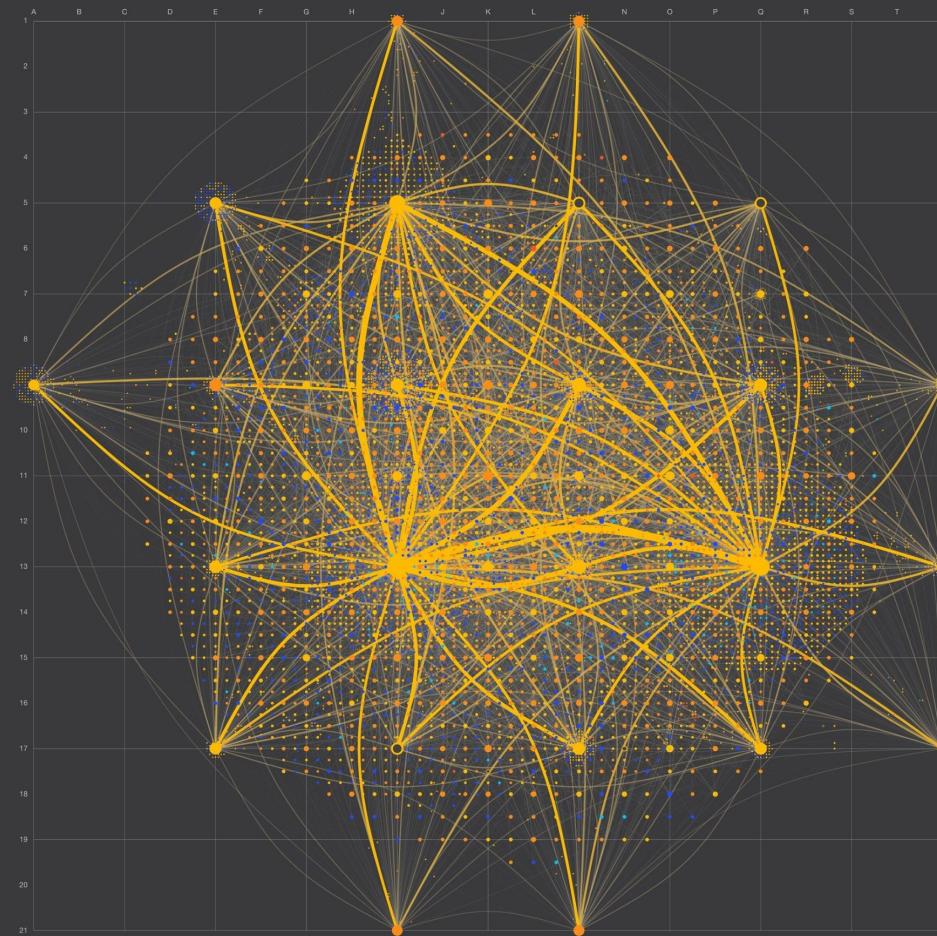
KEGG database access

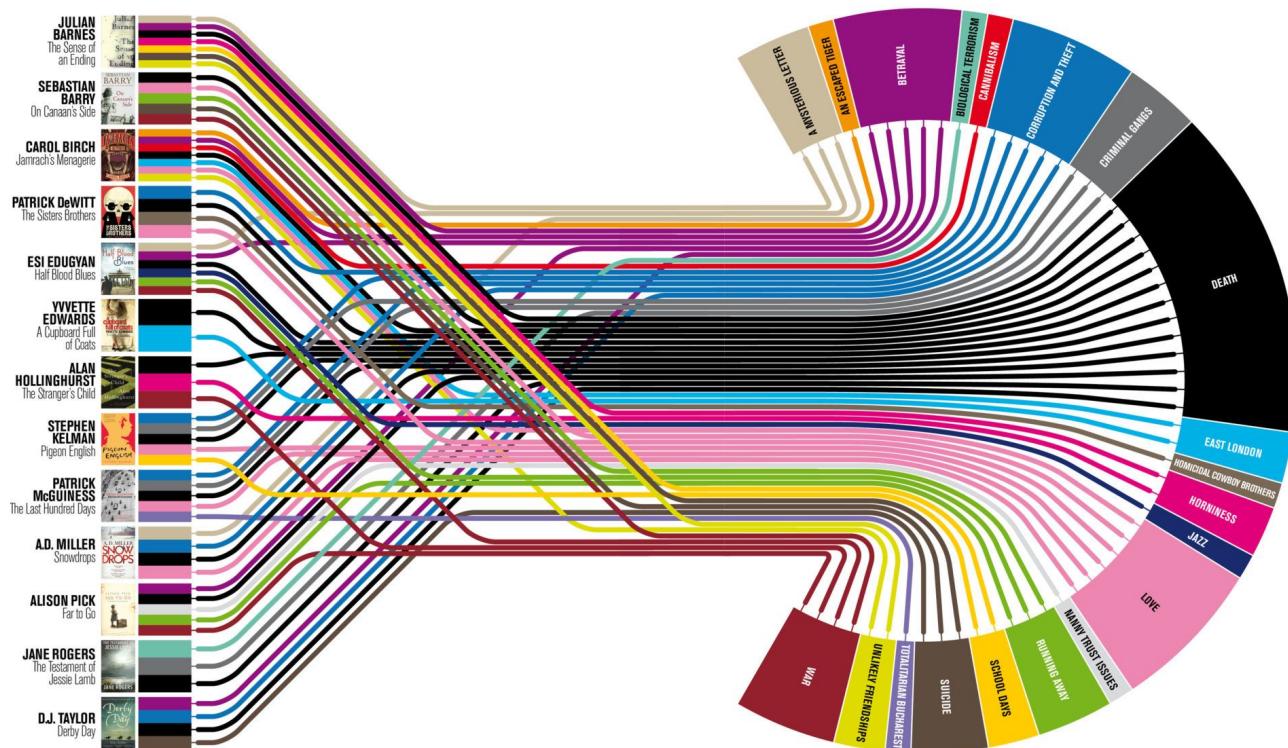


2

The Internet

Topology of Autonomous Systems, 2011.01.02



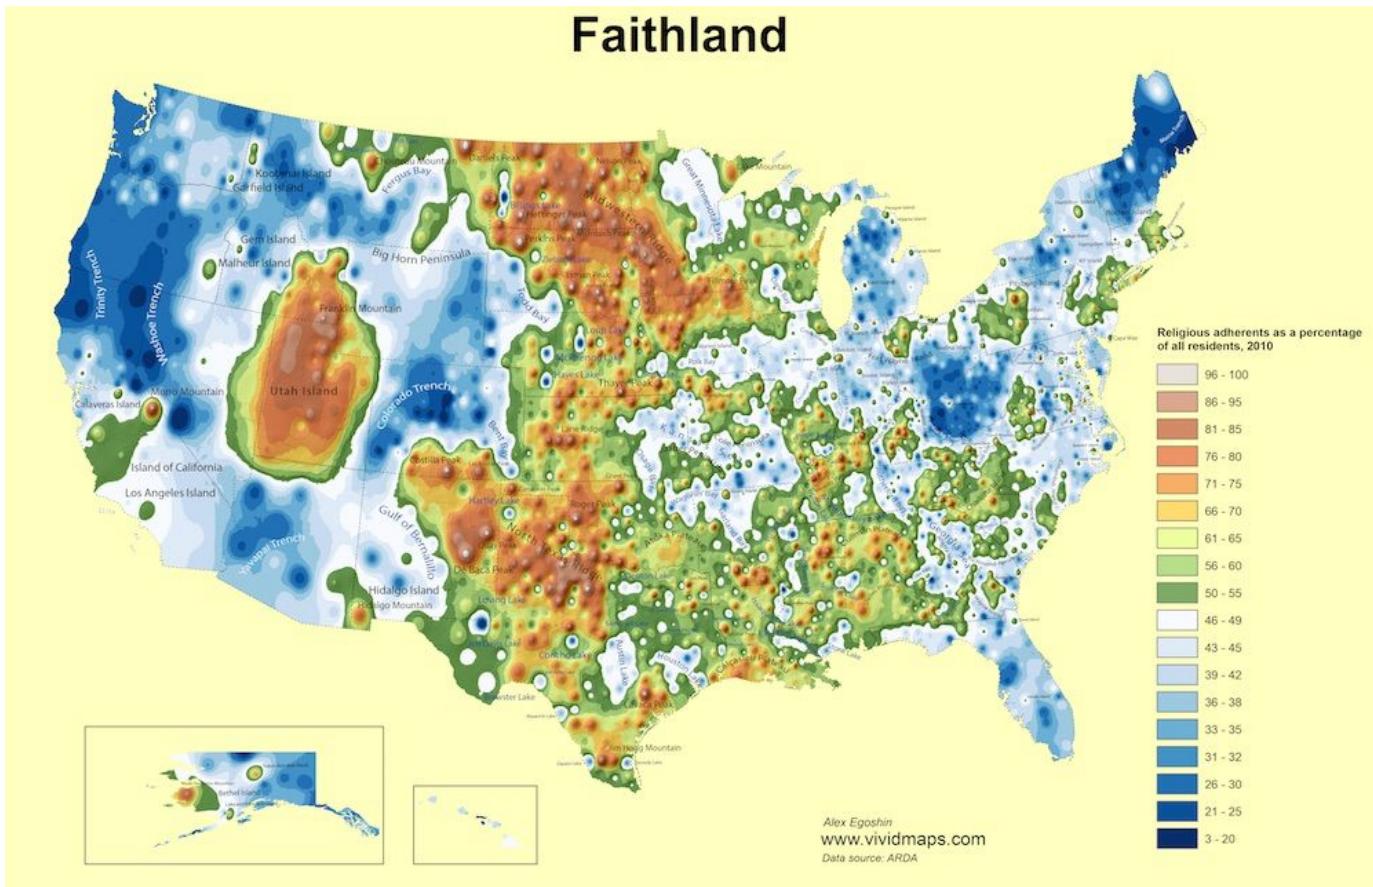


Plot lines

What makes a prize-winning novel? As Julian Barnes wins the Booker Prize, Delayed Gratification's Johanna Kamradt charts the themes of this year's longlisters.

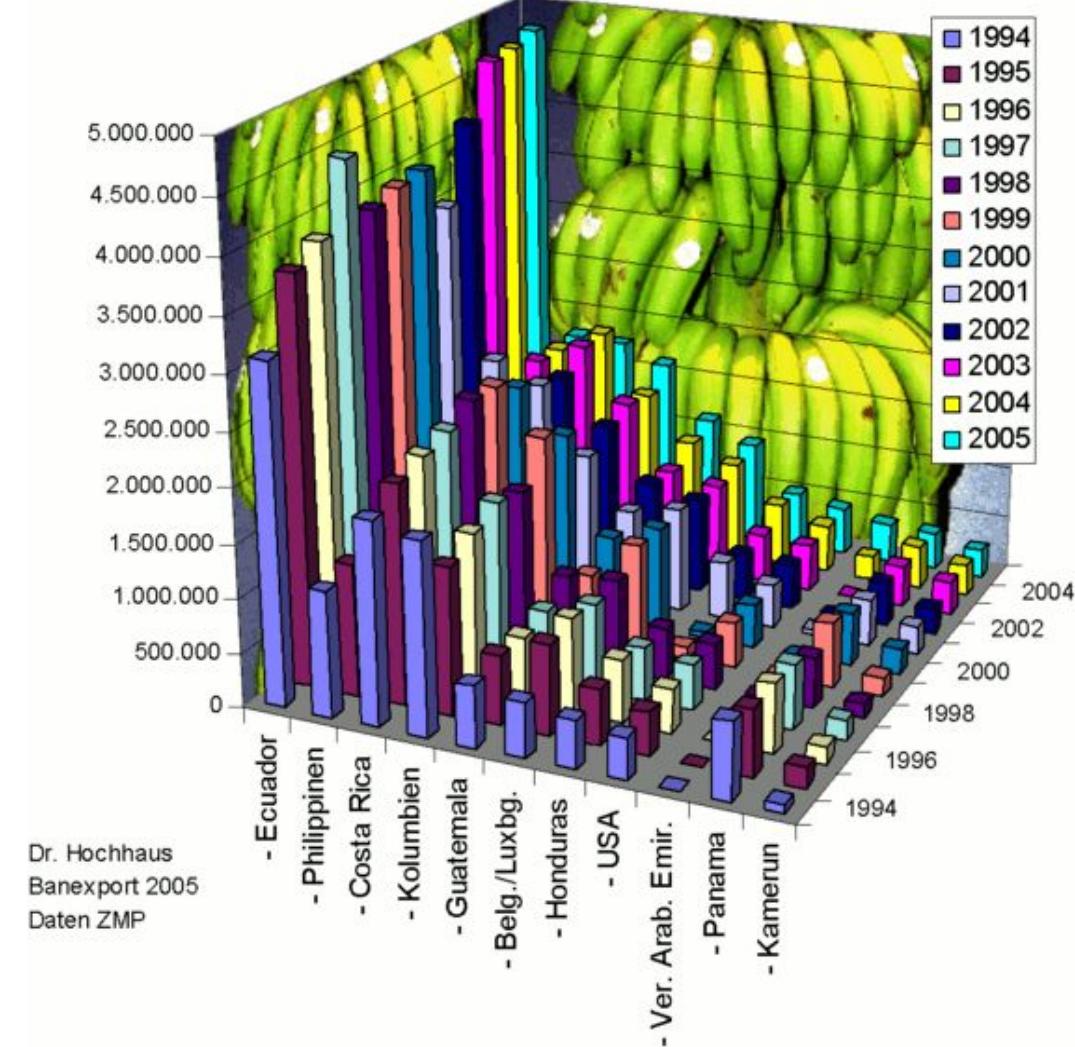
Illustration: Christian Tate





CONFIDENTIAL: Oxmedica Ltd. 10347756

Export von Bananen in Tonnen von 1994-2005

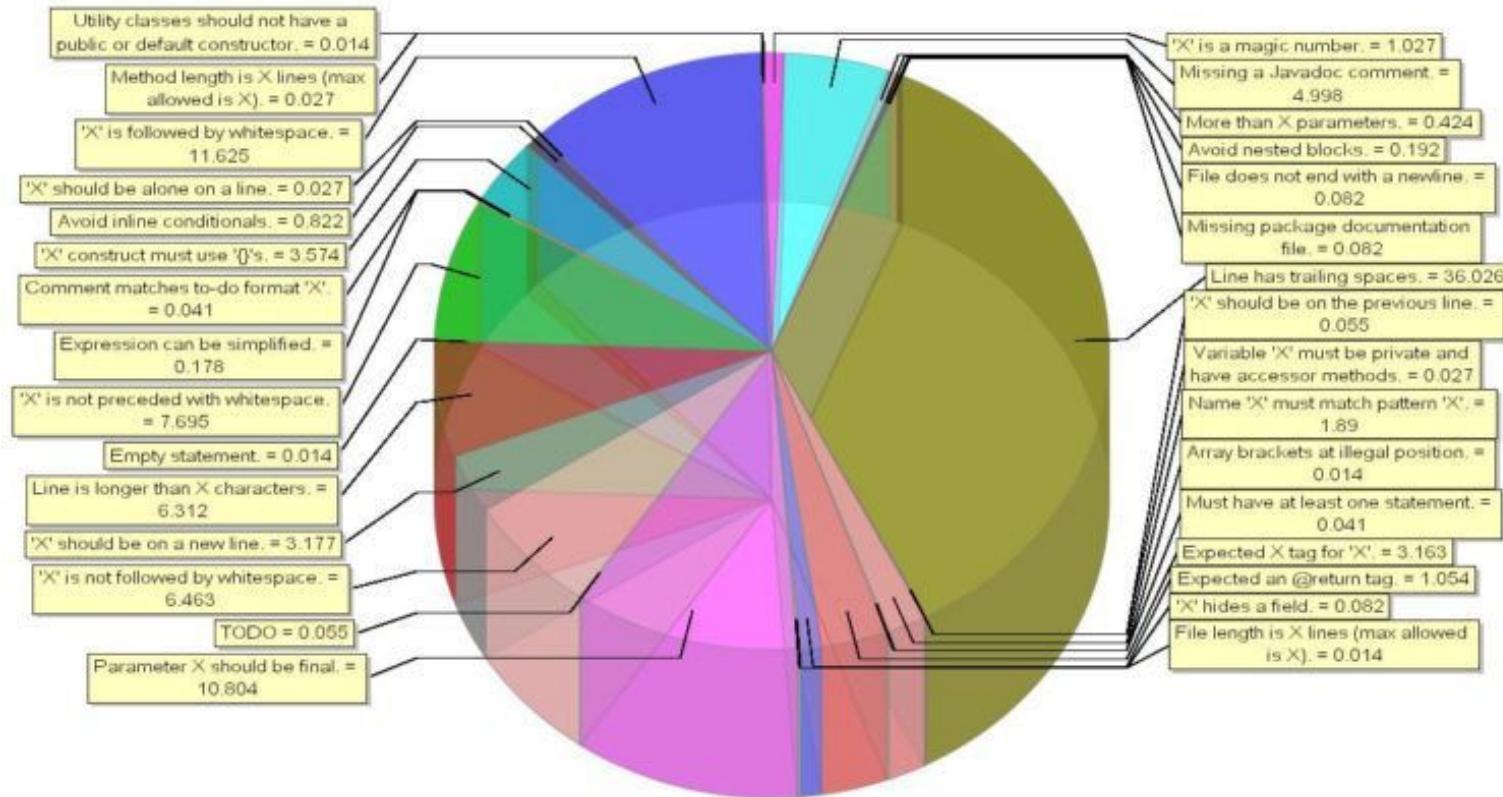


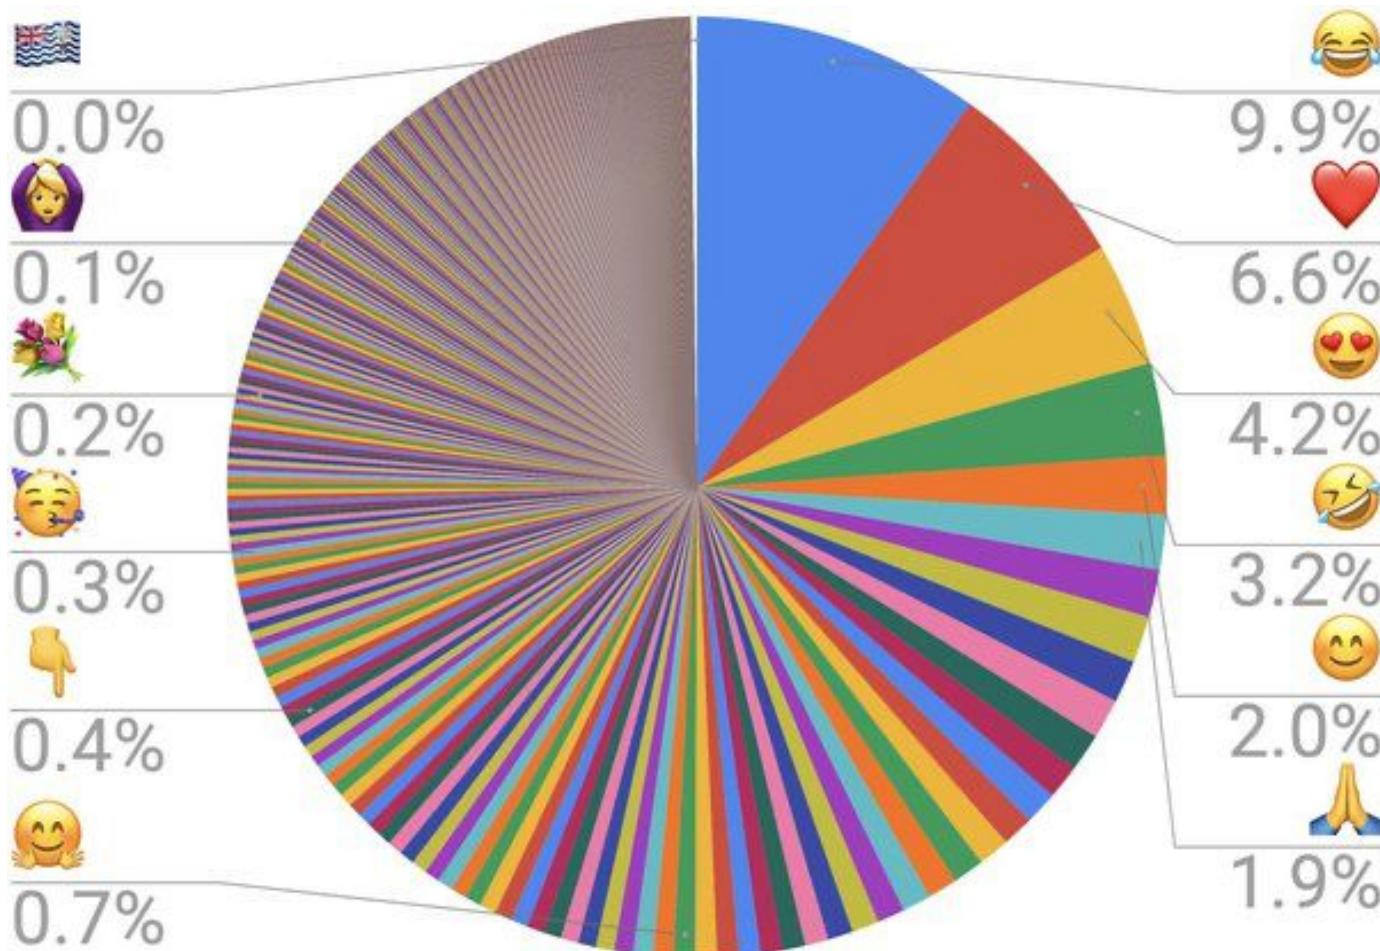
Dr. Hochhaus
Banlexport 2005
Daten ZMP

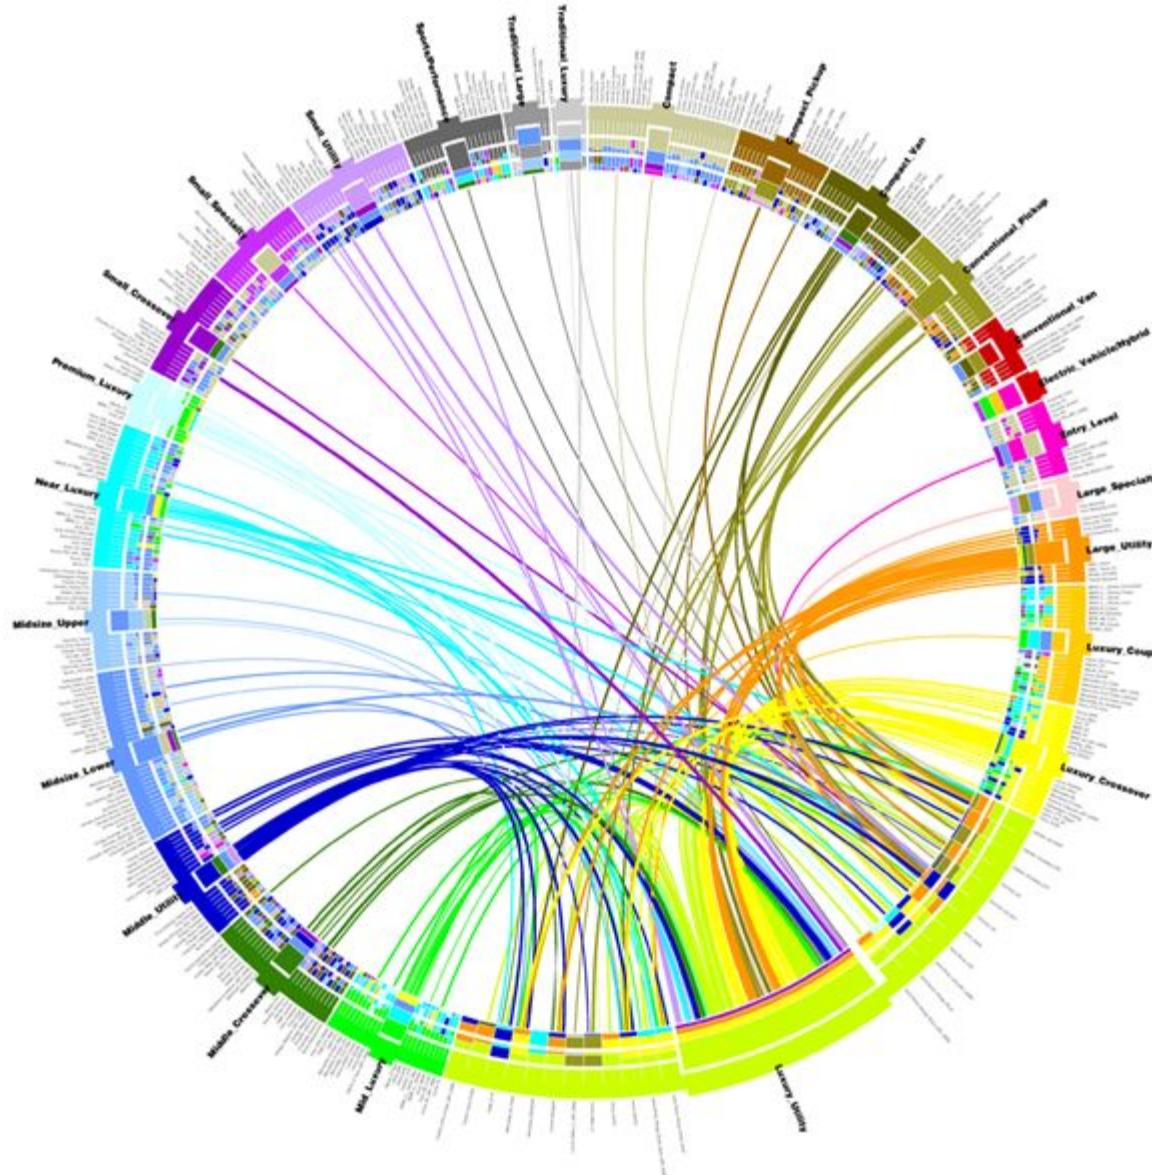


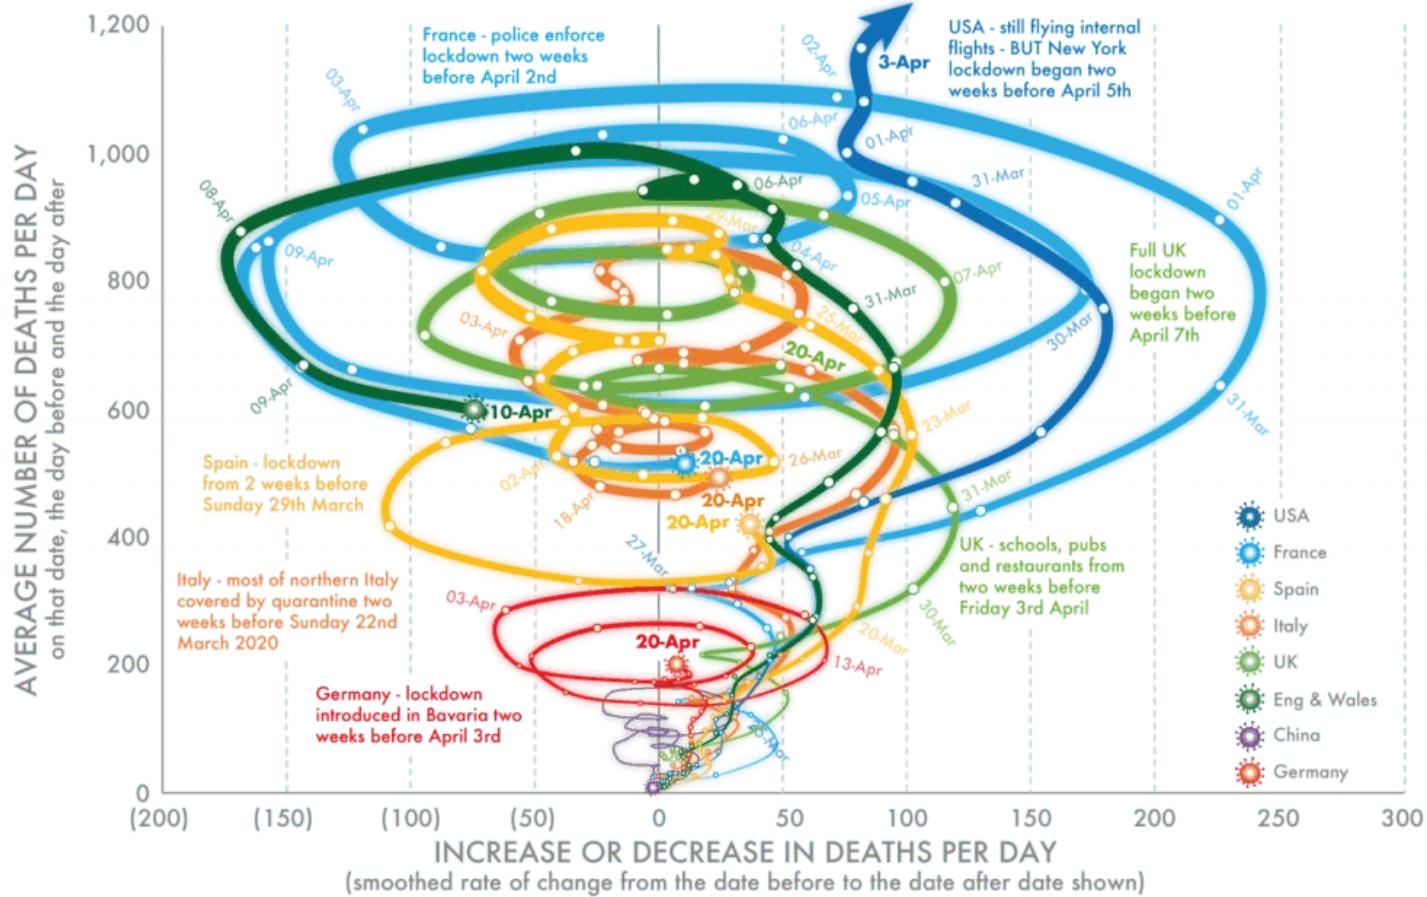
CONFIDENTIAL: Oxmedica Ltd. 10347756











Computational Graphs



CONFIDENTIAL: Oxmedica Ltd. 10347756



⚠ A little warning ⚠



CONFIDENTIAL: Oxmedica Ltd. 10347756



Linear Regression as a Computational Graph



The very simple linear regression we know and love

$$\hat{y}^{(i)} = \hat{\alpha} + \hat{\beta}_1 x_1^{(i)} + \hat{\beta}_2 x_2^{(i)} + \hat{\beta}_3 x_3^{(i)} + \hat{\beta}_4 x_4^{(i)}$$



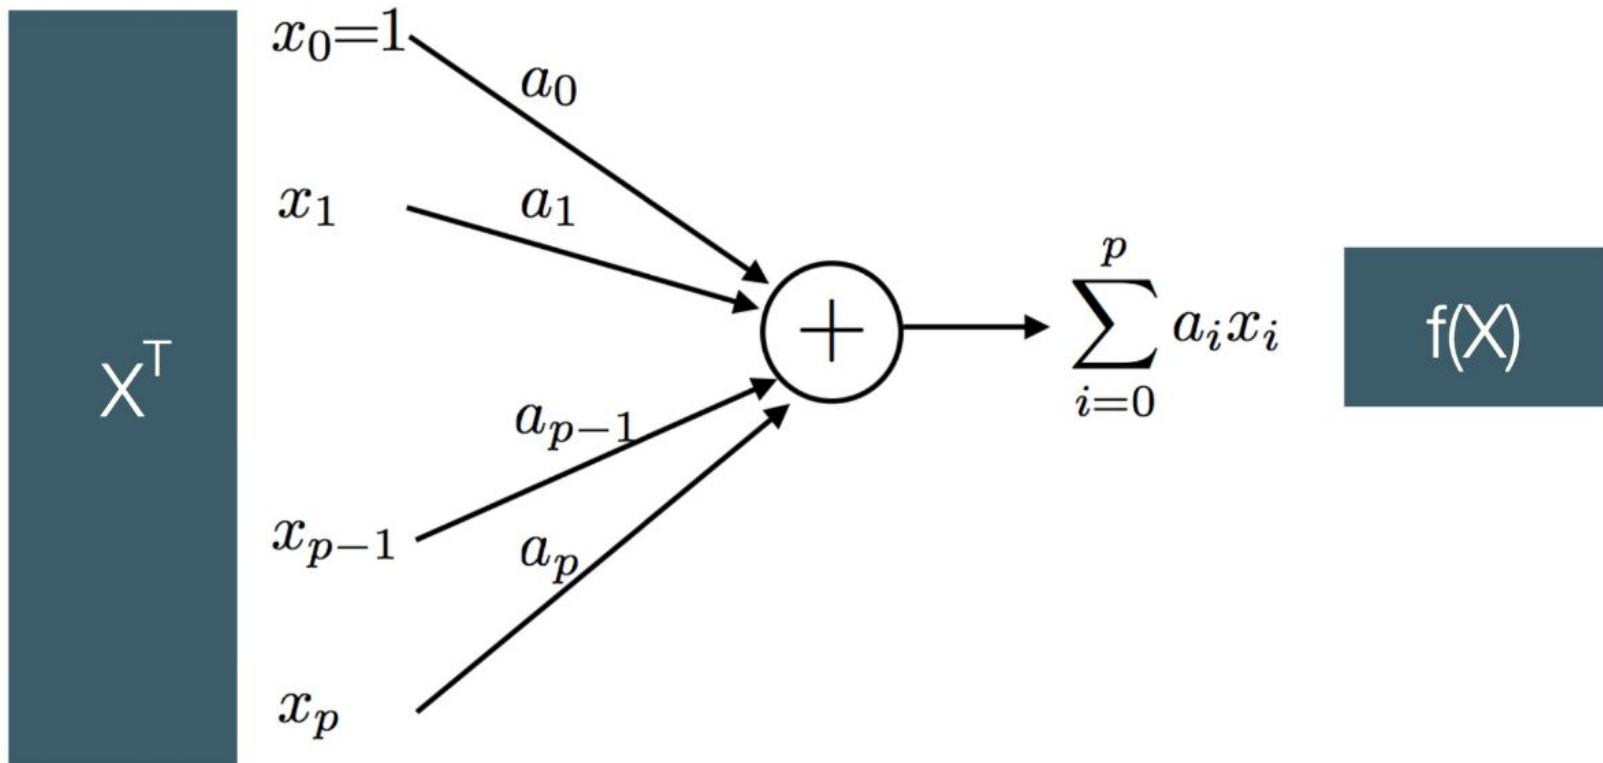
$$\hat{y}^{(i)} = \hat{\alpha} + \hat{\beta}_1 x_1^{(i)} + \hat{\beta}_2 x_2^{(i)} + \hat{\beta}_3 x_3^{(i)} + \hat{\beta}_4 x_4^{(i)}$$

- What are the parameters?
- What type of output does this model provide?



CONFIDENTIAL: Oxmedica Ltd. 10347756

Visualising this as a graph



Activation functions

Whiteboard

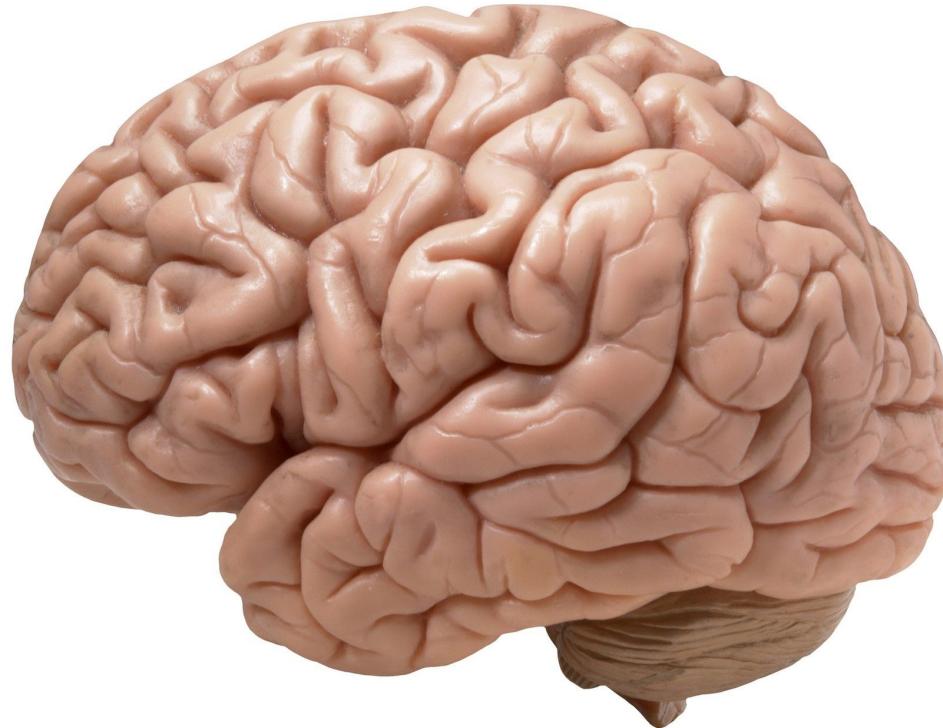


CONFIDENTIAL: Oxmedica Ltd. 10347756

An introduction to neural networks



The brain



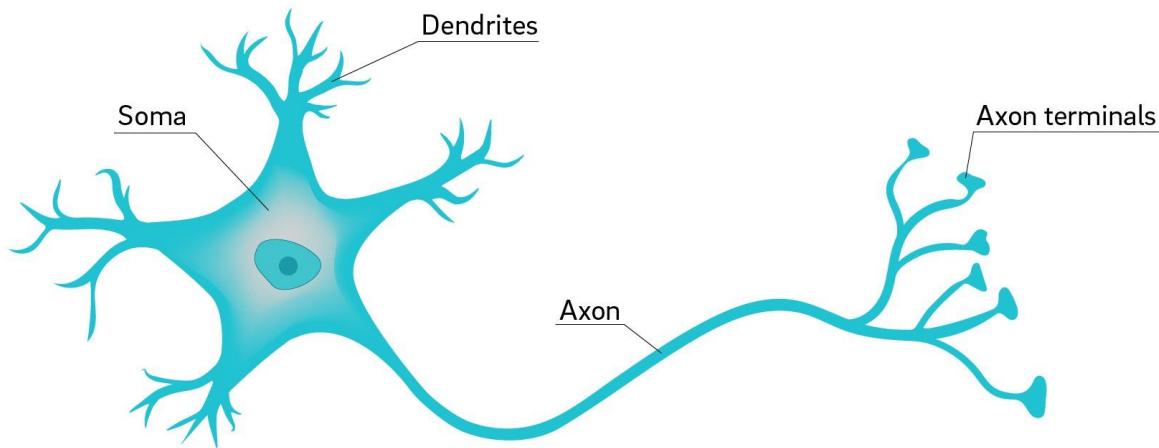
[Source](#)



CONFIDENTIAL: Oxmedica Ltd. 10347756

Biological neurons

Neuron

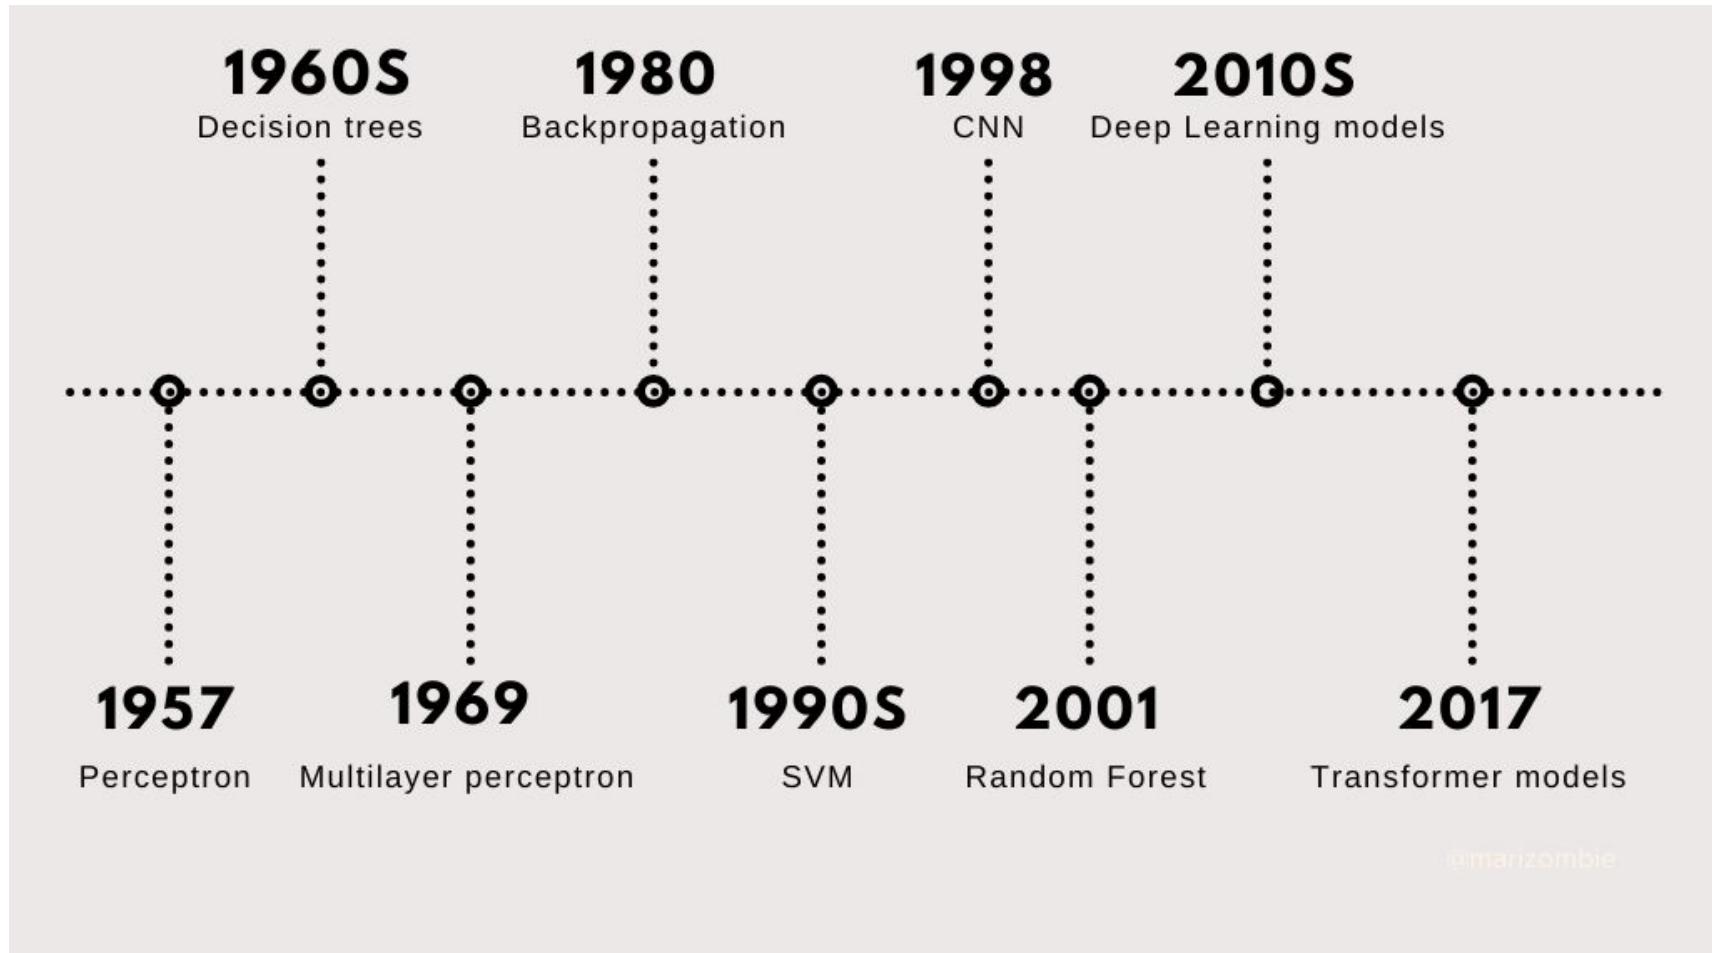


[Source](#)



CONFIDENTIAL: Oxmedica Ltd. 10347756

History of trying to replicate this



[Source](#)



مؤسسة الملك عبد العزيز ورجاله للموهبة والإبداع
King Abdullah Bin Abdulaziz Foundation for Global Initiatives

CONFIDENTIAL: Oxmedica Ltd. 10347756



Artificial neurons

Whiteboard



CONFIDENTIAL: Oxmedica Ltd. 10347756

From one neuron to two...

Whiteboard



From two neurons to many neurons...

Whiteboard

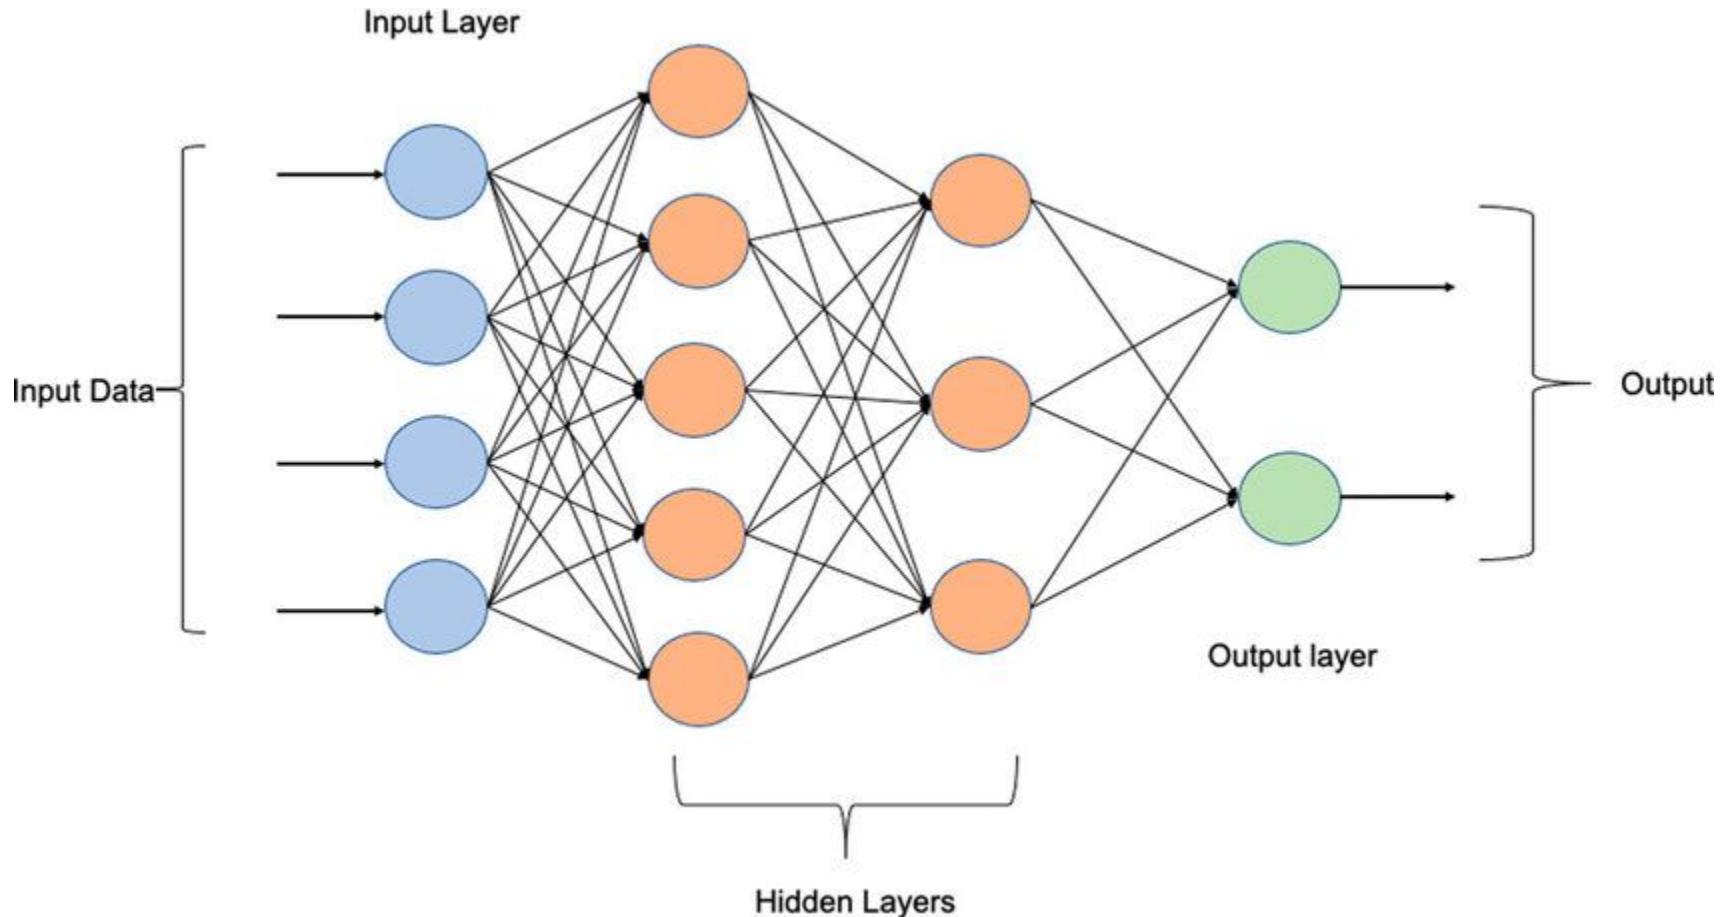


From one layer to two...

Whiteboard



Turning this into larger networks



[Source](#)

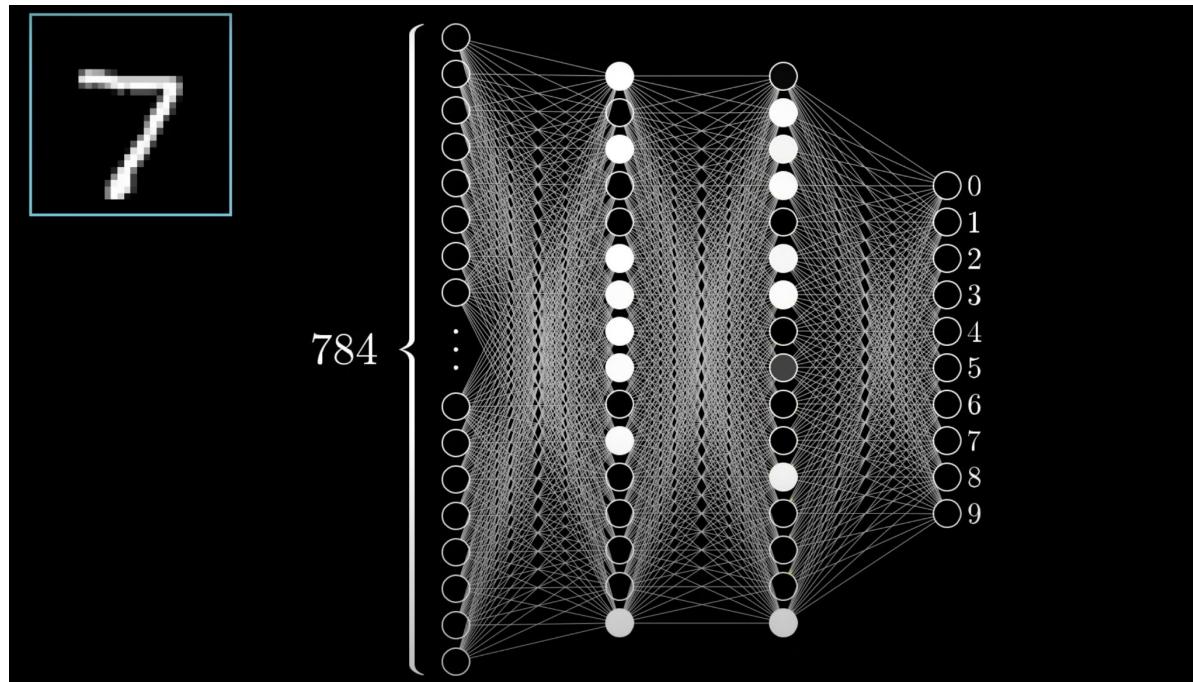
A explanation for MNIST

[Video Comprehension]



An excellent explainer for the MNIST dataset

- Worth making notes on this!
- We'll discuss key questions afterwards
- <https://www.youtube.com/watch?v=aircAruvnKk>



CONFIDENTIAL: Oxmedica Ltd. 10347756

Quick Quiz



مohna

مؤسسة الملك عبد العزيز ورجاله للموهبة والإبداع
King Abdulaziz & Sons Research Foundation for Medicine & Dentistry

CONFIDENTIAL: Oxmedica Ltd. 10347756



Biological: Bigger is (approximately) better

TASK: Guess the average weights of all of these brains
(in kgs)

Elephant: ???

Humans: ???

Monkeys: ???

Hampster: ???



مدون

CONFIDENTIAL: Oxmedica Ltd. 10347756



Biological: Bigger is (approximately) better

TASK: Guess the average weights of all of these brains

Elephant: 5.71kg

Humans: 1.32kg

Monkeys: 0.115kg

Hampster: 0.001kg



Artificial: Bigger is better (*all else equal*)

Task: Guess the number of parameters in these models

GPT-1: ???

GPT-2: ???

GPT-3: ???

GPT-4: ???



Artificial: Bigger is better (*all else equal*)

Task: Guess the number of parameters in these models (in millions/billions/trillions)

GPT-1: 117 million

GPT-2: 1.5 billion

GPT-3: 175 billion

GPT-4: (unknown) 175 bn → >1 trillion



Scaling laws paper (Kaplan et al. 2020)

Scaling Laws for Neural Language Models

Jared Kaplan *

Johns Hopkins University, OpenAI

jaredk@jhu.edu

Sam McCandlish*

OpenAI

sam@openai.com

Tom Henighan

OpenAI

henighan@openai.com

Tom B. Brown

OpenAI

tom@openai.com

Benjamin Chess

OpenAI

bchess@openai.com

Rewon Child

OpenAI

rewon@openai.com

Scott Gray

OpenAI

scott@openai.com

Alec Radford

OpenAI

alec@openai.com

Jeffrey Wu

OpenAI

jeffwu@openai.com

Dario Amodei

OpenAI

damodei@openai.com

<https://arxiv.org/pdf/2001.08361>



Recap questions

1. What are deep learning models? How are they different from other machine learning models?
2. How can a linear regression be represented as a computation graph?
3. What is an activation function?
4. Draw a fully-connected neural network with 4 input features, one hidden layer with 6 neurons and 2 output neurons.
5. Explain the term ‘scaling’ in deep learning.

